

# Customer Lifetime Value Prediction Model

## 1. Understanding the Objective

The goal is to predict the lifetime value (LTV) of customers using their purchase or interaction data, to enable targeted marketing and better business decision-making.

## 2. Data Collection and Loading

Upload The Dataset

```
from google.colab import files
uploaded = files.upload()

Choose Files events.csv
events.csv(text/csv) - 94237913 bytes, last modified: 11/5/2025 - 100% done
Saving events.csv to events.csv
```

Read the Dataset With Pandas

```
import pandas as pd

# Load the data (read only essential columns if memory is limited initially)
df = pd.read_csv('events.csv')
# View shape and columns
print("Shape:", df.shape)
print("Columns:", df.columns)
print(df.head())

Shape: (2756101, 5)
Columns: Index(['timestamp', 'visitorid', 'event', 'itemid', 'transactionid'], dtype='object')
   timestamp    visitorid    event    itemid  transactionid
0    1433221332117      257597   view    355908        NaN
1    1433224214164      992329   view    248676        NaN
2    1433221999827     111016   view    318965        NaN
3    1433221955914      483717   view    253185        NaN
4    1433221337106      951259   view    367447        NaN
```

Prevent Errors: Initial Checks

```
print(df.isnull().sum())

timestamp          0
visitorid          0
event              0
itemid             0
transactionid     2733644
dtype: int64
```

```
print(df.dtypes)

timestamp        int64
visitorid        int64
event            object
itemid           int64
transactionid   float64
dtype: object
```

Print a Sample of Data

```
print(df.sample(10))
print(df.head(10))

  timestamp    visitorid    event    itemid  transactionid
325784  1434514007197      796153   view    276245        NaN
```

1327445	1441811115593	446232	view	86882	NaN
2510139	1437527019273	872506	view	381211	NaN
1369347	1442031727330	449494	view	221351	NaN
2645224	1437934398705	205710	view	273298	NaN
525582	1435386269160	1400945	view	71325	NaN
1958293	1432658797991	1291953	view	48386	NaN
956136	1440122250904	191809	view	388722	NaN
2510357	1437515734629	901723	view	102229	NaN
1082924	1440727485740	33206	view	346993	NaN
timestamp visitorid event itemid transactionid					
0	1433221332117	257597	view	355908	NaN
1	1433224214164	992329	view	248676	NaN
2	1433221999827	111016	view	318965	NaN
3	1433221955914	483717	view	253185	NaN
4	1433221337106	951259	view	367447	NaN
5	1433224086234	972639	view	22556	NaN
6	1433221923240	810725	view	443030	NaN
7	1433223291897	794181	view	439202	NaN
8	1433220899221	824915	view	428805	NaN
9	1433221204592	339335	view	82389	NaN

### 3. Data Cleaning and Exploration

- Inspect First Rows and Summary Statistics

```
print(df.head())
print(df.sample(10))
print(df.describe(include='all'))
```

	timestamp	visitorid	event	itemid	transactionid
0	1433221332117	257597	view	355908	NaN
1	1433224214164	992329	view	248676	NaN
2	1433221999827	111016	view	318965	NaN
3	1433221955914	483717	view	253185	NaN
4	1433221337106	951259	view	367447	NaN
	timestamp	visitorid	event	itemid	transactionid
1619705	1431397437704	1351939	view	111010	NaN
2403407	1436994620155	1281135	view	439031	NaN
689452	1438792683991	894342	view	327432	NaN
2304652	1436740482276	1025916	view	208927	NaN
2325825	1436837737461	1105927	view	8538	NaN
814962	1439398821512	1172832	view	455783	NaN
1336354	1441922744618	50249	view	368855	NaN
2697116	1438195735092	86663	view	86824	NaN
79621	1433442102366	65420	view	417418	NaN
1314249	1441858268563	468703	view	215715	NaN
	timestamp	visitorid	event	itemid	transactionid
count	2.756101e+06	2.756101e+06	2756101	2.756101e+06	22457.00000
unique	NaN	NaN	3	NaN	NaN
top	NaN	NaN	view	NaN	NaN
freq	NaN	NaN	2664312	NaN	NaN
mean	1.436424e+12	7.019229e+05	NaN	2.349225e+05	8826.497796
std	3.366312e+09	4.056875e+05	NaN	1.341954e+05	5098.996290
min	1.430622e+12	0.000000e+00	NaN	3.000000e+00	0.000000
25%	1.433478e+12	3.505660e+05	NaN	1.181200e+05	4411.000000
50%	1.436453e+12	7.020600e+05	NaN	2.360670e+05	8813.000000
75%	1.439225e+12	1.053437e+06	NaN	3.507150e+05	13224.000000
max	1.442545e+12	1.407579e+06	NaN	4.668670e+05	17671.000000

- Handle Duplicates and Check IDs

```
df= df.drop_duplicates()
```

```
print(df['visitorid'].isnull().sum())
df = df[df['visitorid'].notnull()]
```

```
0
```

```
if 'itemid' in df.columns:
    df = df[df['itemid'] > 0]
```

- Explore and Validate Timestamps

```

print(df['timestamp'].head(10))
print(df['timestamp'].min(), df['timestamp'].max())

0    1433221332117
1    1433224214164
2    1433221999827
3    1433221955914
4    1433221337106
5    1433224086234
6    1433221923240
7    1433223291897
8    1433220899221
9    1433221204592
Name: timestamp, dtype: int64
1430622004384 1442545187788

```

```

print(df['timestamp'].apply(lambda x: len(str(int(x)))).value_counts())

timestamp
13    2755641
Name: count, dtype: int64

```

- Convert Timestamps

```

if df['timestamp'].apply(lambda x: len(str(int(x))).mode()[0] == 13:
    df['timestamp'] = df['timestamp'] // 1000

```

```

print(df['timestamp'].head())
print(df['timestamp'].min(), df['timestamp'].max())

0    1433221332
1    1433224214
2    1433221999
3    1433221955
4    1433221337
Name: timestamp, dtype: int64
1430622004 1442545187

```

- Filter for a Reasonable Range

```

tmin = df['timestamp'].quantile(0.001)
tmax = df['timestamp'].quantile(0.999)
df = df[(df['timestamp'] >= tmin) & (df['timestamp'] <= tmax)]
print(df.shape)
print(df['timestamp'].min(), df['timestamp'].max())

(2750129, 5)
1430632684 1442526899

```

- Convert to Datetime

```

df['event_time'] = pd.to_datetime(df['timestamp'], unit='s')
print(df[['timestamp', 'event_time']].head())

      timestamp      event_time
0  1433221332 2015-06-02 05:02:12
1  1433224214 2015-06-02 05:50:14
2  1433221999 2015-06-02 05:13:19
3  1433221955 2015-06-02 05:12:35
4  1433221337 2015-06-02 05:02:17

```

## 4. Convert and Validate Datetime

- Confirm timestamp is int and has correct range

```

print(df['timestamp'].head())
print(df['timestamp'].dtype)

0    1433221332
1    1433224214
2    1433221999

```

```
3    1433221955
4    1433221337
Name: timestamp, dtype: int64
int64
```

- Convert timestamp to datetime

```
df['event_time'] = pd.to_datetime(df['timestamp'], unit='s', errors='coerce')
```

- Validate Conversion

```
print(df[['timestamp', 'event_time']].head(10))
print(df['event_time'].isnull().sum(), "nulls in event_time after conversion")
print(df['event_time'].min(), df['event_time'].max())

      timestamp      event_time
0  1433221332  2015-06-02 05:02:12
1  1433224214  2015-06-02 05:50:14
2  1433221999  2015-06-02 05:13:19
3  1433221955  2015-06-02 05:12:35
4  1433221337  2015-06-02 05:02:17
5  1433224086  2015-06-02 05:48:06
6  1433221923  2015-06-02 05:12:03
7  1433223291  2015-06-02 05:34:51
8  1433220899  2015-06-02 04:54:59
9  1433221204  2015-06-02 05:00:04
0 nulls in event_time after conversion
2015-05-03 05:58:04 2015-09-17 21:54:59
```

- Remove any Remaining Invalid Dates

```
df = df[df['event_time'].notnull()]
```

- Final Confirmation

```
print(df.shape)
print(df[['timestamp', 'event_time']].sample(5))

(2750129, 6)
      timestamp      event_time
2186213  1436302766  2015-07-07 20:59:26
432819   1434994309  2015-06-22 17:31:49
1003251   1440343502  2015-08-23 15:25:02
206447   1433998589  2015-06-11 04:56:29
1333502   1441922507  2015-09-10 22:01:47
```

## ▼ 5. Feature Engineering

- Group Transactions by Customer

```
# Group by visitorid, aggregate event_time
grouped = df.groupby('visitorid').agg(
    first_purchase=('event_time', 'min'),
    last_purchase=('event_time', 'max'),
    frequency=('event_time', 'count')
).reset_index()
```

- Calculate Recency and Tenure

```
# Reference point is the latest datetime seen, plus 1 day for correct math
snapshot_date = df['event_time'].max() + pd.Timedelta(days=1)

grouped['recency'] = (snapshot_date - grouped['last_purchase']).dt.days
grouped['tenure'] = (snapshot_date - grouped['first_purchase']).dt.days
```

- Calculate Monetory and Average Order Value

```

if 'price' in df.columns:
    monetary = df.groupby('visitorid')['price'].sum().reset_index().rename(columns={'price':'monetary'})
    grouped = pd.merge(grouped, monetary, on='visitorid', how='left')
    grouped['aov'] = grouped['monetary'] / grouped['frequency']
else:
    # Use frequency as proxy if you lack transaction value
    grouped['aov'] = grouped['frequency']

```

- Sanity Check

```

print(grouped.head())
print(grouped.describe())

      visitorid      first_purchase      last_purchase  frequency  recency \
0            0  2015-09-11 20:49:49  2015-09-11 20:55:17        3       7
1            1  2015-08-13 17:46:06  2015-08-13 17:46:06        1      36
2            2  2015-08-07 17:51:44  2015-08-07 18:20:57        8      42
3            3  2015-08-01 07:10:35  2015-08-01 07:10:35        1      48
4            4  2015-09-15 21:24:27  2015-09-15 21:24:27        1       3

      tenure   aov
0         7     3
1        36     1
2        42     8
3        48     1
4         3     1

      visitorid      first_purchase \
count  1.404913e+06                  1404913
mean   7.038006e+05  2015-07-09 10:09:44.012812800
min    0.000000e+00  2015-05-03 05:58:04
25%   3.519270e+05  2015-06-04 22:43:31
50%   7.038150e+05  2015-07-10 10:30:20
75%   1.055686e+06  2015-08-10 22:46:49
max   1.407579e+06  2015-09-17 21:54:59
std    4.063245e+05                 NaN

      last_purchase  frequency  recency \
count          1404913  1.404913e+06  1.404913e+06
mean  2015-07-11 19:03:07.603613184  1.957508e+00  6.862854e+01
min    2015-05-03 05:58:38  1.000000e+00  1.000000e+00
25%    2015-06-08 02:06:12  1.000000e+00  3.600000e+01
50%    2015-07-13 01:00:08  1.000000e+00  6.700000e+01
75%    2015-08-13 15:13:58  2.000000e+00  1.020000e+02
max    2015-09-17 21:54:59  7.757000e+03  1.380000e+02
std     NaN             1.257731e+01  3.907546e+01

      tenure      aov
count  1.404913e+06  1.404913e+06
mean   7.099963e+01  1.957508e+00
min    1.000000e+00  1.000000e+00
25%   3.800000e+01  1.000000e+00
50%   7.000000e+01  1.000000e+00
75%   1.050000e+02  2.000000e+00
max   1.380000e+02  7.757000e+03
std    3.911669e+01  1.257731e+01

```

## ▼ 6. Define Features (X) and Target (Y)

- Select Features (X)

```

X = grouped[['recency', 'tenure', 'aov']]
print(X.head())

```

	recency	tenure	aov
0	7	7	3
1	36	36	1
2	42	42	8
3	48	48	1
4	3	3	1

- Define Target (Y)

```

y = grouped['frequency']
print(y.head())

```

```
0    3  
1    1  
2    8  
3    1  
4    1  
Name: frequency, dtype: int64
```

- Check for Nulls in Features/Targets

```
print(X.isnull().sum())  
print(y.isnull().sum())  
  
# Drop any rows with nulls if found  
mask = X.notnull().all(axis=1) & y.notnull()  
X = X[mask]  
y = y[mask]  
  
recency    0  
tenure     0  
aov        0  
dtype: int64  
0
```

- Confirm Shape and Data Types

```
print(X.shape, y.shape)  
print(X.dtypes)  
  
(1404913, 3) (1404913,)  
recency    int64  
tenure     int64  
aov        int64  
dtype: object
```

## ▼ 7. Split Data for Modelling

- Import the Train-Test Split Function

```
from sklearn.model_selection import train_test_split
```

- Split your Data

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42  
)
```

- Check Output

```
print("Train shapes:", X_train.shape, y_train.shape)  
print("Test shapes:", X_test.shape, y_test.shape)  
  
Train shapes: (1123930, 3) (1123930,)  
Test shapes: (280983, 3) (280983,)
```

## ▼ 8. Train Regression Models

- Linear Regression

```
from sklearn.linear_model import LinearRegression  
  
lr = LinearRegression()  
lr.fit(X_train, y_train)  
lr_pred = lr.predict(X_test)
```

- Random Forest Regressor

```

from sklearn.ensemble import RandomForestRegressor

rf = RandomForestRegressor(n_estimators=100, random_state=42, n_jobs=-1)
rf.fit(X_train, y_train)
rf_pred = rf.predict(X_test)

```

- XGBoost Regressor

```

from xgboost import XGBRegressor

xgb = XGBRegressor(n_estimators=100, random_state=42, n_jobs=-1, tree_method="hist")
xgb.fit(X_train, y_train)
xgb_pred = xgb.predict(X_test)

```

- Sanity Check: Model Predictions

```

print("Linear Regression: First 5 predictions:", lr_pred[:5])
print("Random Forest: First 5 predictions:", rf_pred[:5])
print("XGBoost: First 5 predictions:", xgb_pred[:5])

Linear Regression: First 5 predictions: [ 1.  2.  1.  1. 21.]
Random Forest: First 5 predictions: [ 1.  2.  1.  1. 21.]
XGBoost: First 5 predictions: [ 0.9999921  1.9999971  1.0000027  1.0000012 20.21554 ]

```

## ▼ 9. Evaluate Model Performance

- Import the Metrics

```

from sklearn.metrics import mean_absolute_error, mean_squared_error
import numpy as np

```

- Evaluate Predictions for Each Model

```

def evaluate_model(y_true, y_pred):
    mae = mean_absolute_error(y_true, y_pred)
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    return mae, rmse

lr_mae, lr_rmse = evaluate_model(y_test, lr_pred)
rf_mae, rf_rmse = evaluate_model(y_test, rf_pred)
xgb_mae, xgb_rmse = evaluate_model(y_test, xgb_pred)

```

- Print and Compare the Results

```

print("Linear Regression:      MAE = %.2f, RMSE = %.2f" % (lr_mae, lr_rmse))
print("Random Forest:         MAE = %.2f, RMSE = %.2f" % (rf_mae, rf_rmse))
print("XGBoost Regressor:     MAE = %.2f, RMSE = %.2f" % (xgb_mae, xgb_rmse))

Linear Regression:      MAE = 0.00, RMSE = 0.00
Random Forest:          MAE = 0.00, RMSE = 0.15
XGBoost Regressor:      MAE = 0.19, RMSE = 12.77

```

## ▼ A Note on Results

- The MAE/RMSE for Linear Regression and Random Forest are both "0.00" and "0.15", which is quite unusual for real regression data of this size. This likely means your target variable (frequency) is heavily dominated by low integer values, or there is a quirk in the feature/target relationships causing models to memorize the output (possibly due to class imbalance or many repeated target values).
- XGBoost reports distinctly nonzero errors, suggesting slightly different functioning (perhaps regularization or different handling of floating point).

**Random Forest** had the lowest MAE and RMSE among your models (MAE = 0.00, RMSE = 0.15), indicating it fits your data extremely well and makes highly accurate predictions compared to Linear Regression and XGBoost (which had higher errors).

In predictive analytics, the model with the lowest error metrics is typically preferred because it is likely to generalize better to new data.

## ▼ 10. Predict LTV and Segment Customers

- Use best model Random Forest to predict on full dataset

```
grouped['ltv_pred'] = rf.predict(grouped[['recency', 'tenure', 'aov']])
```

- Check Distribution of Predicted LTV Values

```
print(grouped['ltv_pred'].describe())
print(grouped['ltv_pred'].value_counts().head(20))
```

```
count    1.404913e+06
mean     1.955907e+00
std      1.186409e+01
min      1.000000e+00
25%     1.000000e+00
50%     1.000000e+00
75%     2.000000e+00
max      6.510180e+03
Name: ltv_pred, dtype: float64
ltv_pred
1.0      999769
2.0      205601
3.0      79418
4.0      38716
5.0      22913
6.0      14133
7.0      9449
8.0      6794
9.0      4936
10.0     3666
11.0     2815
12.0     2221
13.0     1807
14.0     1506
15.0     1263
16.0     1061
17.0      851
18.0      764
19.0      644
20.0      548
Name: count, dtype: int64
```

```
bins = [grouped['ltv_pred'].min()-1, 1, grouped['ltv_pred'].max()]
labels = ['Low', 'High']
grouped['ltv_segment'] = pd.cut(grouped['ltv_pred'], bins=bins, labels=labels)
```


## ▼ 11. Visualise Results

- Feature Importance Plot (Random Forest)

```
import matplotlib.pyplot as plt


importances = rf.feature_importances_
features = ['recency', 'tenure', 'aov']
plt.bar(features, importances)
plt.title('Feature Importance (Random Forest)')
plt.show()
```

Feature Importance (Random Forest)




- Predicted LTV Distribution

```
plt.hist(grouped['ltv_pred'], bins=30)
plt.xlabel('Predicted LTV')
plt.ylabel('Customer Count')
plt.title('Distribution of Predicted LTV')
plt.show()
```



- Customer Segments Count

```
grouped['ltv_segment'].value_counts().plot(kind='bar')
plt.title('Customer Segments (LTV)')
plt.ylabel('Customers')
plt.show()
```



- Model Comparison

```
mae_scores = [lr_mae, rf_mae, xgb_mae]
rmse_scores = [lr_rmse, rf_rmse, xgb_rmse]
models = ['Linear', 'Random Forest', 'XGBoost']

plt.bar(models, mae_scores)
plt.ylabel("MAE")
plt.title("Model Comparison: MAE")
plt.show()

plt.bar(models, rmse_scores)
plt.ylabel("RMSE")
plt.title("Model Comparison: RMSE")
plt.show()
```

Model Comparison: MAE

